

HW3 Hadoop Mapreduce

ID: r08921a07

NAME: 曾梓豪

1. Run the application in a Hadoop Docker container

the choose to use the docker container from the link:

<https://github.com/sdwangntu/hadoop-cluster>

First, I used command “swarm init” then created network with command “docker network create -d overlay --attachable my-attachable-network”

```
how123480@how123480-computer ~/hadoop-cluster master ? docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
4761f335bb4e        bridge              bridge              local
7a8a51f0f6c3        docker_gwbridge     bridge              local
4b87a8ad194c        host                host                local
6mc8af0c9lbz        ingress             overlay             swarm
hvq1n08ce2t4        my-attachable-network overlay             swarm
29c21b58e44c        none                null                local
how123480@how123480-computer ~/hadoop-cluster master ?
```

Second, I used commands below to launch the cluster:

1. docker run --hostname=mysql --name mysql --network my-attachable-network -d sdwangntu/hive-metastore-db
2. docker run --hostname=hadoop-master --name hadoop-master --network my-attachable-network -d sdwangntu/hadoop3hbase-spark-hive
3. docker run --hostname=hadoop-worker --name hadoop-worker --network my-attachable-network -d sdwangntu/hadoop3hbase-spark-hive

```
how123480@how123480-computer ~/hadoop-cluster master ? docker ps
CONTAINER ID        IMAGE                                     NAMES                COMMAND              CREATED      STATUS      PORTS
0cd4d4d44710       sdwangntu/hive-metastore-db            mysql                "docker-entrypoint.s..." 3 days ago  Up 2 minutes  3306/tcp, 33060/tcp
f2ad2842c76a       sdwangntu/hadoop3hbase-spark-hive      hadoop-worker        "bash /start-hadoop..." 3 days ago  Up 2 minutes  2181/tcp, 8042/tcp, 8088/tcp, 8888/tcp, 9864/tcp, 9870/tcp
e0e403e72859       sdwangntu/hadoop3hbase-spark-hive      hadoop-worker        "bash /start-hadoop..." 3 days ago  Up 2 minutes  2181/tcp, 8042/tcp, 8088/tcp, 8888/tcp, 9864/tcp, 9870/tcp
p, 16000/tcp, 16010/tcp, 19888/tcp  hadoop-master
how123480@how123480-computer ~/hadoop-cluster master ?
```

Third, I deployed the development container with command “docker run --hostname=hadoop-dev --name hadoop-dev -v \$(pwd):/home --network my-attachable-network -d sdwangntu/hadoop3hbase-spark-hive”

```
how123480@how123480-computer ~/hadoop-cluster master ? docker ps
CONTAINER ID        IMAGE                                     NAMES                COMMAND              CREATED      STATUS      PORTS
5b4ae0cf9373       sdwangntu/hadoop3hbase-spark-hive      hadoop-dev           "bash /start-hadoop..." 3 days ago  Up 6 seconds  2181/tcp, 8042/tcp, 8088/tcp, 8888/tcp, 9864/tcp, 9870/tcp, 16000/tcp, 16010/tcp, 19888/tcp
how123480@how123480-computer ~/hadoop-cluster master ?
```

Fourth, use command “docker exec -it {container ID} bash” to get the shell of the container. then used command “printenv” to print the environment variable. We can know hadoop in the directory “/opt/hadoop”

```

root@hadoop-dev:~# printenv
LD_LIBRARY_PATH=/opt/hadoop/lib/native:
LS_COLORS=rs=0:di=01;34:ln=01;36:mh=00:pi=40;33:so=01;35:do=01;35:bd=40;33;01:cd=40;33;01:or=40;33;01:mi=00:su=37;41:sg=30;43:ca=30;41:tw=30;42:ow=34;42:st=37;44:ex=01;32:*.tar=0
1;31:*.tgz=01;31:*.arc=01;31:*.arj=01;31:*.taz=01;31:*.lha=01;31:*.lz4=01;31:*.lzh=01;31:*.lzma=01;31:*.tlz=01;31:*.txz=01;31:*.tzo=01;31:*.t7z=01;31:*.zip=01;31:*.z=01;31:*.Z=01
;31:*.dz=01;31:*.gz=01;31:*.lrz=01;31:*.lzo=01;31:*.xz=01;31:*.zst=01;31:*.tztst=01;31:*.bz2=01;31:*.bz=01;31:*.tbz=01;31:*.tbz2=01;31:*.tz=01;31:*.deb=01;31:*.rpm=01;3
1:*.jar=01;31:*.war=01;31:*.ear=01;31:*.sar=01;31:*.rar=01;31:*.alz=01;31:*.ace=01;31:*.zoo=01;31:*.cpio=01;31:*.7z=01;31:*.rz=01;31:*.cab=01;31:*.wlm=01;31:*.swm=01;31:*.dwm=01;
31:*.esd=01;31:*.jpg=01;35:*.jpeg=01;35:*.mjpeg=01;35:*.gif=01;35:*.bmp=01;35:*.pbm=01;35:*.pgm=01;35:*.ppm=01;35:*.tga=01;35:*.xbm=01;35:*.xpm=01;35:*.tif=01;35:*.t
iff=01;35:*.png=01;35:*.svg=01;35:*.svgz=01;35:*.mng=01;35:*.pcx=01;35:*.mov=01;35:*.mpg=01;35:*.mp2=01;35:*.m4v=01;35:*.mkv=01;35:*.webm=01;35:*.ogm=01;35:*.mp3=00;36:*.mpc=
35:*.mp4v=01;35:*.vob=01;35:*.qt=01;35:*.nuv=01;35:*.wmv=01;35:*.asf=01;35:*.ram=01;35:*.rmvb=01;35:*.flc=01;35:*.avi=01;35:*.fli=01;35:*.flv=01;35:*.gl=01;35:*.dl=01;35:*.xcf=01;
35:*.xwd=01;35:*.yuv=01;35:*.cgm=01;35:*.cfl=01;35:*.ogv=01;35:*.ogx=01;35:*.aac=00;36:*.au=00;36:*.flac=00;36:*.m4a=00;36:*.mid=00;36:*.midi=00;36:*.mka=00;36:*.mp3=00;36:*.mpc=
00;36:*.ogg=00;36:*.ra=00;36:*.wav=00;36:*.oga=00;36:*.opus=00;36:*.spx=00;36:*.xspf=00;36:
LESSCLOSE=/usr/bin/lesspipe %s %s
HADOOP_HOME=/opt/hadoop
HOSTNAME=hadoop-dev
HIVE_HOME=/opt/hive
JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
PWD=/
HOME=/root
SPARK_HOME=/opt/spark
TERM=xterm
HBASE_HOME=/opt/hbase
SHLVL=1
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/opt/hadoop/bin:/opt/hbase/bin:/opt/spark/bin:/opt/hive/bin
LESSOPEN=| /usr/bin/lesspipe %s
_/usr/bin/printenv
root@hadoop-dev:~#

```

edit the “mapper” code of lecture example, “wordcount” to parse the log file.

```

1 #!/usr/bin/env python
2 import sys
3 import re
4
5 month_map={
6     "Jan": "01",
7     "Feb": "02",
8     "Mar": "03",
9     "Apr": "04",
10    "May": "05",
11    "Jun": "06",
12    "Jul": "07",
13    "Aug": "08",
14    "Sep": "09",
15    "Oct": "10",
16    "Nov": "11",
17    "Dec": "12"
18 }
19 def month_sub(time):
20     for key in month_map:
21         if(key in time):
22             return time.replace(key,month_map[key])
23     return time
24
25 for line in sys.stdin:
26     try:
27         time = re.split('\[|\]',line)[1]
28     except:
29         continue
30
31     time = ":".join(time.split(":", 2)[:2])
32     time = month_sub(time)
33
34     time = re.split('\[/\|:',time)
35     time = "{}-{}-{} T {}:00:00.000".format(time[2],time[0],time[1],time[3])
36     print '%s\t%s' % (time,1)
37
~

```

then use hadoop-streaming to start the program.

```

1 hdfs dfs -rm -r -f log_outdir
2
3 hadoop \
4     jar "/opt/hadoop/share/hadoop/tools/lib/hadoop-streaming-3.1.2.jar" \
5     -D mapred.map.tasks=6 \
6     -mapper "$PWD/mapper.py" \
7     -reducer "$PWD/reducer.py" \
8     -input "log" \
9     -output "log_outdir" \
10    -file "$PWD/mapper.py" \
11    -file "$PWD/reducer.py"
~
~

```

Finally, I get the result of log analysis.

```
Peak Map Physical memory (bytes)=314232832
Peak Map Virtual memory (bytes)=2659741696
Peak Reduce Physical memory (bytes)=191184896
Peak Reduce Virtual memory (bytes)=2660491264
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=193383
File Output Format Counters
  Bytes Written=2825
2020-10-29 14:38:23,413 INFO streaming.StreamJob: Output directory: log_outdir
root@hadoop-dev:~/log# hdfs dfs -ls log_outdir
Found 2 items
-rw-r--r--  1 root supergroup          0 2020-10-29 14:38 log_outdir/_SUCCESS
-rw-r--r--  1 root supergroup    2825 2020-10-29 14:38 log_outdir/part-00000
root@hadoop-dev:~/log# _
```

```
root@hadoop-dev:~/log# hdfs dfs -cat log_outdir/part-00000
2004-07-03 T 16:00:00.000      27
2004-07-03 T 17:00:00.000      25
2004-07-03 T 18:00:00.000      24
2004-07-03 T 19:00:00.000      26
2004-07-03 T 20:00:00.000      20
2004-07-03 T 21:00:00.000      23
2004-07-03 T 22:00:00.000      29
2004-07-03 T 23:00:00.000      22
2004-08-03 T 00:00:00.000      21
2004-08-03 T 01:00:00.000      21
2004-08-03 T 02:00:00.000      27
2004-08-03 T 03:00:00.000      22
2004-08-03 T 04:00:00.000      26
2004-08-03 T 05:00:00.000      37
2004-08-03 T 06:00:00.000      17
2004-08-03 T 07:00:00.000      31
2004-08-03 T 08:00:00.000      44
2004-08-03 T 09:00:00.000      63
2004-08-03 T 10:00:00.000      39
2004-08-03 T 11:00:00.000      34
2004-08-03 T 12:00:00.000      45
2004-08-03 T 13:00:00.000      37
2004-08-03 T 14:00:00.000      23
2004-08-03 T 15:00:00.000       9
2004-08-03 T 16:00:00.000       2
2004-08-03 T 17:00:00.000       2
```

2. Write the java version of part(a)

I also edited the “mapper” code in lecture example “wordcount”, and following is the code of “mapper” section code.

In order to run the java, I set the environment variable in “/etc/profile” then “source” it.

```
export CLASSPATH=.:$HADOOP_HOME/share/hadoop/common/hadoop-common-3.1.2.jar:$HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-client-core-3.1.2.jar:$HADOOP_HOME/share/hadoop/common/lib/commons-cli-1.2.jar:$CLASSPATH
root@hadoop-dev:~# _
```

```

7 import org.apache.hadoop.mapreduce.Reducer;
8 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat; import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
9
10
11 public class Log {
12
13     public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable>{
14         String month_sub(String time){
15             Map<String, String> month_map = new HashMap<String, String>() {{
16                 put("Jan", "01");
17                 put("Feb", "02");
18                 put("Mar", "03");
19                 put("Apr", "04");
20                 put("May", "05");
21                 put("Jun", "06");
22                 put("Jul", "07");
23                 put("Aug", "08");
24                 put("Sep", "09");
25                 put("Oct", "10");
26                 put("Nov", "11");
27                 put("Dec", "12");}};
28
29             for (Iterator it = month_map.entrySet().iterator(); it.hasNext();){
30                 Map.Entry mapEntry = (Map.Entry) it.next();
31
32                 if(time.contains((String)mapEntry.getKey())){
33                     time = time.replace((String)mapEntry.getKey(), (String)mapEntry.getValue());
34                     return time;
35                 }
36             }
37             return time;
38         }
39
40         private final static IntWritable one = new IntWritable(1);
41         private Text word = new Text();
42         public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
43
44             //parsing start
45             String time = value.toString().split("\\[\\W\\]")[1];
46             String[] tokens = time.split(":");
47             time = tokens[0] + ":" + tokens[1];
48             time = month_sub(time);
49             tokens = time.split("\\[\\W\\]");
50             time = String.format("%s-%s-%s T %s:00:00.000",tokens[2],tokens[0],tokens[1],tokens[3]);
51             //parsing end
52             word.set(time);
53             context.write(word, one);
54         }
55     }
56
57     public static class IntSumReducer extends Reducer<Text,IntWritable,Text,IntWritable> {

```

28,0-1

20%

Then I use the following script to start the hadoop program and display the result.

```

1 source /etc/profile
2 hdfs dfs -rm -r -f log-outdir
3 rm -rf ./build
4 mkdir build
5 javac -d ./build Log.java
6 cd build
7 jar cvf Log.jar *
8 yarn jar Log.jar Log log log-outdir
9 cd ..
10 hdfs dfs -cat log-outdir/part-r-00000
~

```

Following is the result of java version

```

Virtual memory (bytes) snapshots:531762224
Total committed heap usage (bytes)=522715136
Peak Map Physical memory (bytes)=295581824
Peak Map Virtual memory (bytes)=2655596544
Peak Reduce Physical memory (bytes)=189452288
Peak Reduce Virtual memory (bytes)=2662215680

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=172903

File Output Format Counters
Bytes Written=2825
2004-07-03 T 16:00:00.000 27
2004-07-03 T 17:00:00.000 25
2004-07-03 T 18:00:00.000 24
2004-07-03 T 19:00:00.000 26
2004-07-03 T 20:00:00.000 20
2004-07-03 T 21:00:00.000 23
2004-07-03 T 22:00:00.000 29
2004-07-03 T 23:00:00.000 22
2004-08-03 T 00:00:00.000 21
2004-08-03 T 01:00:00.000 21
2004-08-03 T 02:00:00.000 27
2004-08-03 T 03:00:00.000 22
2004-08-03 T 04:00:00.000 26
2004-08-03 T 05:00:00.000 37
2004-08-03 T 06:00:00.000 17
2004-08-03 T 07:00:00.000 31
2004-08-03 T 08:00:00.000 44
2004-08-03 T 09:00:00.000 63

```

3. Source code and running script github link:

<https://github.com/how123480/simple-hadoop>